

1) Consider dataset $\{x_1, \dots, x_n\}$. $J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|_2^2$.

a) what is minimum value of J when $K = n$?

→ suppose each datapoint aligns with the prototype of some cluster.

Then each prototype has 1 datapoint in it such that

· For each k , $r_{nk}=1$ for only one datapoint x_n .

· For that $r_{nk}=1$, $\|x_n - \mu_k\|_2^2 = 0$.

Then, $J = 0$.

Since $J \geq 0$, it follows that the minimum of J is $\boxed{J=0}$.

(sum of squared terms)

b) consider the regularized objective function $J = \sum_{k=1}^K [\lambda \|\mu_k\|_2^2 + \sum_{n=1}^N r_{nk} \|x_n - \mu_k\|_2^2]$.

Suppose all r_{nk} are known. Find $\arg \min_{\mu_k} (\lambda \|\mu_k\|_2^2 + \sum_{n=1}^N r_{nk} \|x_n - \mu_k\|_2^2)$.

$$\frac{\partial J}{\partial \mu_k} = \frac{\partial}{\partial \mu_k} \left(\lambda \|\mu_k\|_2^2 + \sum_{n=1}^N r_{nk} \|x_n - \mu_k\|_2^2 \right)$$

$$= 2\lambda \mu_k + \sum_{n=1}^N -2r_{nk} (x_n - \mu_k)$$

$$= 0$$

$$\Rightarrow 2 \left(\lambda \mu_k - \sum_{n=1}^N r_{nk} (x_n - \mu_k) \right) = 0$$

$$\Rightarrow \lambda \mu_k = \sum_{n=1}^N r_{nk} x_n - \sum_{n=1}^N r_{nk} \mu_k$$

$$\Rightarrow \mu_k (\lambda + \sum_{n=1}^N r_{nk}) = \sum_{n=1}^N r_{nk} x_n$$

$$\Rightarrow \mu_k = \frac{\sum_{n=1}^N r_{nk} x_n}{\lambda + \sum_{n=1}^N r_{nk}}$$




- Zach Bejan

The optimal μ_k is $\mu_k = \frac{\sum_{n=1}^N r_{nk} x_n}{\lambda + \sum_{n=1}^N r_{nk}}$


This regularization would affect the refitting step of K-means because it directly affects how μ_k is calculated.

Namely, as λ gets large, μ_k approaches 0.

Consequently, regularization causes all of the cluster centers to concentrate toward each other at the origin.



unregularized



regularized

2) Unlabeled data $x_n \in \mathbb{R}^m$, $n=1, \dots, N$. Use L_1 distance; $J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|_1$.

Let $C_n = \{k \mid r_{nk} = 1\}$. Goal: minimize $f(\mu_k) = \sum_{n \in C_n} \|x_n - \mu_k\|_1$ wrt μ_k .

Let x_{ni} be i^{th} element in x_n , μ_{ki} be i^{th} element in μ_k .

$$\text{Then, } f(\mu_k) = \sum_{n \in C_n} \|x_n - \mu_k\|_1 = \sum_{n \in C_n} \sum_{i=1}^m |x_{ni} - \mu_{ki}| = \sum_{i=1}^m \sum_{n \in C_n} |x_{ni} - \mu_{ki}|.$$

Hence, we can optimize for each element of μ_k separately.

\hookrightarrow Note: $\sum_{n \in C_n}$ iterates over all data points in k^{th} cluster.

- - -

a) Find $\partial f(x)$ of $f(x) = |x|$ for $x < 0$, $x > 0$ and $x = 0$.

When $x < 0$, $f(x)$ is differentiable with $f'(x) = -1$.

Hence the subgradient, which equals the derivative of f at points where f is differentiable, is given by $\partial f(x) = \{-1\}$.


Similarly when $x > 0$, $\partial f(x) = \{1\}$.

At $x = 0$, use the definition of subgradient to write that $\forall z \in \text{dom}f$, g satisfies $f(z) \geq f(0) + g(z-0) \Leftrightarrow \underline{|z|} \geq g \geq \underline{+z}$.

$|z| \geq g \geq +z$ is satisfied iff subgradient $g \in [-1, 1]$.

Hence for $x = 0$, $\partial f(x) = [-1, 1]$.

Graphical representation:



$$\partial f(x) = \begin{cases} \{-1\}, & x < 0 \\ [-1, 1], & x = 0 \\ \{1\}, & x > 0 \end{cases}$$

b) Show that $\bar{y}^* = y_{\frac{N_k+1}{2}}$ minimizes $\sum_{j=1}^{N_k} |y_j - \bar{y}| = f(\bar{y})$ given the assumptions.

Suppose $f(\bar{y})$ is convex and subdifferentiable everywhere.

Then, f is subdifferentiable at $\bar{y}^* = y_{\frac{N_k+1}{2}}$.

Since N_k is odd, \bar{y}^* exists and is the median of y_1, \dots, y_{N_k} .

Since $y_1 < y_2 < \dots < y_{N_k}$, the following is true:

$$1. (y_j - \bar{y}^*) < 0 \text{ for } j = 1, \dots, \frac{N_k+1}{2} - 1$$

$$2. (y_j - \bar{y}^*) = 0 \text{ for } j = \frac{N_k+1}{2}$$

$$3. (y_j - \bar{y}^*) > 0 \text{ for } j = \frac{N_k+1}{2} + 1, \dots, N_k$$

Then by part a, the subgradient of $|y_j - \bar{y}^*|$ is the following:

$$1. -1 \text{ for } j = 1, \dots, \frac{N_k+1}{2} - 1 \quad 2. [-1, 1] \text{ for } j = \frac{N_k+1}{2}$$

$$3. +1 \text{ for } j = \frac{N_k+1}{2} + 1, \dots, N_k$$

$$\text{Hence, } \partial f(\bar{y}^*) = \sum_{j=1}^{\frac{N_k+1}{2}-1} (-1) + \sum_{j=\frac{N_k+1}{2}+1}^{N_k} (+1) + [\text{subgradient of } |y_{\frac{N_k+1}{2}} - \bar{y}^*|].$$

Yet, there is an equal number of terms on each side of the median because N_k is odd, so the first two sums add to equal 0.

Also, $0 \in \{\text{subgradient of } |y_{\frac{N_k+1}{2}} - \bar{y}^*|\}$.

Hence, $0 \in \partial f(\bar{y}^*)$.

Recalling above that f is convex and subdifferentiable at \bar{y}^* , by the provided theorem, it follows that \bar{y}^* minimizes $\sum_{j=1}^{N_k} |y_j - \bar{y}|$.

□

c) 2 Step algorithm that minimizes J ...

o) Initialization: Start with some random initialization of the prototypes μ_k for $1 \leq k \leq K$.

1) Assignment: For each datapoint, find the prototype that is closest to it. Mathematically, set $r_{nk} = \begin{cases} 1 & \text{if } k = \underset{1 \leq j \leq K}{\operatorname{argmin}} \|x_n - \mu_j\|^2 \\ 0 & \text{else} \end{cases}$

2) Refitting: Suppose all indicators r_{nk} are known.

Must find prototypes that minimize $J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|_1$.

Let $\mu_k = [\text{median of all datapoints in cluster } k]$.

↳ consequence of pt. b.

One iteration includes step 1. and step 2.

Keep going until a stopping criterion.

This algorithm is advantageous compared to k-means because it is computationally more simple in the step of refitting.

Whereas k-means sets μ_k equal to the mean of datapoints in cluster k , a $O(n)$ operation, this algorithm sets μ_k to the median, which if implemented well can be done in $O(1)$ time.

Also note that the median does not consider drastic outliers, whereas the mean does, showing another reason why this algorithm is advantageous.

3) Consider $A = \begin{bmatrix} 3 & 2 \\ 2 & 0 \end{bmatrix}$

a) Find eigenvalues and eigenvectors of A , then normalize eigenvectors.

$$P(\lambda) = \det(A - \lambda I) = \det \begin{bmatrix} 3-\lambda & 2 \\ 2 & -\lambda \end{bmatrix} = (3-\lambda)(-\lambda) - (2)(2) = -3\lambda + \lambda^2 - 4$$

$$\Rightarrow P(\lambda) = (\lambda - 4)(\lambda + 1)$$

$\Rightarrow \lambda_1 = 4, \lambda_2 = -1$ are eigenvalues

Look at respective eigenspaces.

$$\hookrightarrow \lambda_1 = 4: \ker(A - 4I) = \ker \begin{bmatrix} -1 & 2 \\ 2 & -4 \end{bmatrix} = \text{span} \begin{pmatrix} 2 \\ 1 \end{pmatrix}$$

$\Rightarrow \begin{pmatrix} 2 \\ 1 \end{pmatrix}$ is an eigenvector for $\lambda_1 = 4$.

$$\text{Normalized, obtain } u_1 = \frac{\begin{pmatrix} 2 \\ 1 \end{pmatrix}}{\sqrt{2^2+1^2}} = \frac{1}{\sqrt{5}} \begin{pmatrix} 2 \\ 1 \end{pmatrix}$$

$$\lambda_2 = -1: \ker(A + I) = \ker \begin{bmatrix} 4 & 2 \\ 2 & 1 \end{bmatrix} = \text{span} \begin{pmatrix} 1 \\ -2 \end{pmatrix}$$

$\Rightarrow \begin{pmatrix} 1 \\ -2 \end{pmatrix}$ is an eigenvector for $\lambda_2 = -1$.

$$\text{Normalized, obtain } u_2 = \frac{1}{\sqrt{5}} \begin{pmatrix} 1 \\ -2 \end{pmatrix}$$

$\lambda_1 = 4$ is eigenvalue w/ eigenvector $\frac{1}{\sqrt{5}} \begin{pmatrix} 2 \\ 1 \end{pmatrix}$

$\lambda_2 = -1$ is eigenvalue w/ eigenvector $\frac{1}{\sqrt{5}} \begin{pmatrix} 1 \\ -2 \end{pmatrix}$

b) Eigenvalue decomposition of A given by $A = Q\Lambda Q^{-1}$ where Q 's i^{th} column has i^{th} eigenvector and Λ is diagonal with corresponding eigenvalues $\Lambda_{ii} = \lambda_i$.

$$\text{Hence, } A = \begin{bmatrix} \frac{1}{\sqrt{5}} & \frac{3}{\sqrt{5}} \\ -\frac{3}{\sqrt{5}} & \frac{1}{\sqrt{5}} \end{bmatrix} \begin{bmatrix} -1 & 0 \\ 0 & 4 \end{bmatrix} \begin{bmatrix} \frac{1}{\sqrt{5}} & \frac{3}{\sqrt{5}} \\ -\frac{3}{\sqrt{5}} & \frac{1}{\sqrt{5}} \end{bmatrix}^{-1}$$

$$\therefore A = \begin{bmatrix} \frac{1}{\sqrt{5}} & \frac{3}{\sqrt{5}} \\ -\frac{3}{\sqrt{5}} & \frac{1}{\sqrt{5}} \end{bmatrix} \begin{bmatrix} -1 & 0 \\ 0 & 4 \end{bmatrix} \begin{bmatrix} \frac{1}{\sqrt{5}} & \frac{3}{\sqrt{5}} \\ -\frac{3}{\sqrt{5}} & \frac{1}{\sqrt{5}} \end{bmatrix}^{-1}$$

4] A symmetric real matrix M is positive definite if $\forall z \neq 0, z^T M z > 0$.

a) Let $A = \begin{bmatrix} 9 & 6 \\ 6 & a \end{bmatrix}$. what should a satisfy so A is positive definite?

Let $z = \begin{bmatrix} z_1 \\ z_2 \end{bmatrix}$ be some arbitrary vector.

$$\begin{aligned} \text{Then, } z^T A z &= [z_1 \ z_2] \begin{bmatrix} 9 & 6 \\ 6 & a \end{bmatrix} \begin{bmatrix} z_1 \\ z_2 \end{bmatrix} = [z_1 \ z_2] \begin{bmatrix} 9z_1 + 6z_2 \\ 6z_1 + az_2 \end{bmatrix} \\ &= z_1(9z_1 + 6z_2) + z_2(6z_1 + az_2) \\ &= 9z_1^2 + 12z_1z_2 + az_2^2 \\ &= 9z_1^2 + 12z_1z_2 + 4z_2^2 + az_2^2 - 4z_2^2 \\ &= (3z_1 + 2z_2)^2 + az_2^2 - 4z_2^2 \end{aligned}$$

Now, if $z_2 = 0$ then $z_1 \neq 0$ and $(3z_1 + 2z_2)^2 + az_2^2 - 4z_2^2 = 9z_1^2 > 0$.

If $z_2 \neq 0$, then $(3z_1 + 2z_2)^2 + az_2^2 - 4z_2^2 > 0$ provided that

$$az_2^2 - 4z_2^2 > 0 \Leftrightarrow \underline{a > 4}.$$

A is positive definite
if $a > 4$

b) Suppose B is positive definite. Show B^{-1} is positive definite.

Pick some column vector $z \neq 0$.

B^{-1} exists. So, \exists vector w of same dimension as z such that $w = B^{-1}z$.

$w = B^{-1}z \Rightarrow z = Bw$ and $z^T = w^T B^T$.

$$\text{Then, } z^T B^{-1} z = (w^T B^T) B^{-1} (Bw)$$

$$= w^T B^T w$$

$$= (w^T B w)^T$$

$$= w^T B w \quad \text{b/c } w^T B w \text{ is a scalar.}$$

$$> 0 \quad \text{by positive definiteness of } B$$

$\Rightarrow B^{-1}$ is positive definite by definition.

□

c) Show that data covariance matrix S in PCA is positive semi-definite.

S is defined as $S = \frac{1}{N} \sum_{n=1}^N (x_n - \bar{x})(x_n - \bar{x})^T$ of dimension $D \times D$.

$\hookrightarrow N$ is number of data points, $N > 0$.

x_n is a datapoint, $x_n \in \mathbb{R}^D$.

$\bar{x} = \frac{1}{N} \sum_{n=1}^N x_n$ is the mean of all datapoints.

Let v be some arbitrary column vector dimension $D \times 1$ such that $v \neq 0$.

$$\begin{aligned} \text{Then, } v^T S v &= v^T \frac{1}{N} \sum_{n=1}^N (x_n - \bar{x})(x_n - \bar{x})^T v \\ &= \frac{1}{N} \sum_{n=1}^N (v^T x_n - v^T \bar{x})(x_n^T - \bar{x}^T) v \\ &= \frac{1}{N} \sum_{n=1}^N (v^T x_n - v^T \bar{x})(x_n^T v - \bar{x}^T v) \\ &= \frac{1}{N} \sum_{n=1}^N (v^T x_n - v^T \bar{x})^2 \quad \text{b/c } v^T x_n = x_n^T v, \quad v^T \bar{x} = \bar{x}^T v \\ &\geq 0 \end{aligned}$$

Hence, S is positive semidefinite.

□

Zack Berger

Question 5

Part A - Visualization


Using imread and imshow, obtained the following graphic of the Bruin Bear:



Part B - Algorithm

Using the specifications for the K-means algorithm and running 10 iterations yielded the following graph of J vs. iterations.

The graph's asymptotic behavior as the number of iterations gets large indicates that the K-means algorithm will converge after many iterations. The insignificant changes in J as the number of iterations could be used to initiate a stopping condition for the algorithm.



Part C - Compression

As K increases, the quality of the compressed image improves. More of the original colors are retained, and thus the overall resolution of the image is greater...

$K = 4$

- After last iteration, $J = 1.718261597502673 \times 10^8$



$K = 8$

- After last iteration, $J = 8.043382692074148 \times 10^7$



$K = 8$

- After last iteration, $J = 3.742910684965859 \times 10^7$



Part D - Compression Ratio

To store the original image, you would need 2,880,000 bits.

$$300 \times 400 = 120,000 \text{ pixels}$$

Each pixel stores 3 elements vector taking $8 \times 3 = 24$ bits

$$120,000 \text{ (pixels)} \times 24 \text{ (bits/pixel)} = 2,880,000 \text{ bits}$$

To calculate how many bits are required to store the image using K centers, do the following:

$$300 \times 400 = 120,000 \text{ pixels}$$

Each pixel stores 1 index = 8 bits

$$120,000 \text{ (pixels)} \times 8 \text{ (bits/pixel)} = 960,000 \text{ bits}$$

K values stored, each with 24 bit precision

$$K \times 24 \text{ bits} = 24K \text{ bits}$$

$$\text{Total storage} = (24K + 960,000) \text{ bits}$$

This computation lead to the following:

- K = 4: 960,096 bits
- K = 8: 960,192 bits
- K = 16: 960,384 bits

Compute compression ratios as (Uncompressed Size) / (Compressed Size):

- K = 4: 2.9997
- K = 8: 2.9994
- K = 16: 2.9988

Code Attachment

```
% Import 300x400 picture of Bruin Bear, then display it
pic = imread('/Users/zackberger/Desktop/ML/HW/HW_7/UCLA_Bruin.jpg');

% Initialize K-means algorithm; K > 0
K = 16;

prototypes = zeros(3,K);           % Initialize prototypes to zero
prototypes(:,1) = [229, 249, 250]; % Set the first prototype to first pixel

% Set the rest of the prototypes using furthest-first heuristic
for k = 2 : K

    % Start with the first pixel as the furthest pixel
    % fpd = furthest_pixel_distance
    furthest_pixel = [229;249;250];
    fpd = 0;

    % Iterate over all pixels
    for r = 1 : 300
        for c = 1 : 400

            current_pixel = double([pic(r,c,1);pic(r,c,2);pic(r,c,3)]);

            % Find the min distance of current_pixel to previously computed prototypes
            % Denote it cpd. Start with distance to first prototype
            cpd = norm(current_pixel - prototypes(:,1))^2;

            % Iterate over all previously computed prototypes
            for p = 2 : (k - 1)
                % Obtain distance to next prototype (npd)
                npd = norm(current_pixel - prototypes(:,p))^2;
                if npd < cpd
                    cpd = npd;
                end
            end

            % If the closest prototype distance is the furthest pixel
            % distance, we have a new furthest pixel
            if cpd > fpd
                fpd = cpd;
                furthest_pixel = current_pixel;
            end

        end
    end

    % Set the furthest pixel as the next prototype
    prototypes(:,k) = furthest_pixel;

end

% Run K-means for 10 iterations
J = zeros(1,10);                 % Track objective for each iteration
pixel_indices = zeros(300,400);   % Each pixel refers to its closest prototype
Num_Iterations = 10;
```

```

for i = 1 : Num_Iterations

    % Assign each pixel to closest prototype
    for r = 1 : 300
        for c = 1 : 400

            current_pixel = double([pic(r,c,1);pic(r,c,2);pic(r,c,3)]);

            % Calculate closest prototype.
            % Start with distance to first.
            % prototype
            cpd = norm(current_pixel - prototypes(:,1))^2;
            closest_prototype = 1;

            % Iterate over all other prototypes
            for p = 2 : K
                npd = norm(current_pixel - prototypes(:,p))^2;
                if npd < cpd
                    cpd = npd;
                    closest_prototype = p;
                end
            end

            % Set closest prototype
            pixel_indices(r,c) = closest_prototype;

        end
    end

    % Now, pixel_indices contains the number prototype that each pixel is
    % a part of. Re-estimate center of each cluster using this info
    prototypes = zeros(3,K); % Refresh the prototype values
    num_samples = zeros(1,K); % Number of datapoints in each cluster

    for r = 1 : 300
        for c = 1 : 400

            current_pixel = double([pic(r,c,1);pic(r,c,2);pic(r,c,3)]);

            num_samples(1,pixel_indices(r,c)) = num_samples(1,pixel_indices(r,c)) + 1;
            prototypes(:, pixel_indices(r,c)) = prototypes(:, pixel_indices(r,c)) + current_pixel;

        end
    end

    prototypes = prototypes ./ num_samples;

    % Calculate objective function
    total_dist = 0;
    for r = 1 : 300
        for c = 1 : 400

            for k = 1 : K

                current_pixel = double([pic(r,c,1);pic(r,c,2);pic(r,c,3)]);

                if pixel_indices(r,c) == k
                    total_dist = total_dist + norm(current_pixel - prototypes(:,k))^2;
                end

            end

        end
    end

```

```

        end
    end
end

J(l) = total_dist;
end

% Set the picture to correct pixels
for r = 1 : 300
    for c = 1 : 400
        pic(r,c,:) = prototypes(:, pixel_indices(r,c));
    end
end

imshow(pic);

% Plot objective function vs. iterations
% iterations = [1,2,3,4,5,6,7,8,9,10];
% scatter(iterations, J, 'filled');

% title("J Value vs. Iteration");
% xlabel("Iteration Number");
% ylabel("J Value");

```